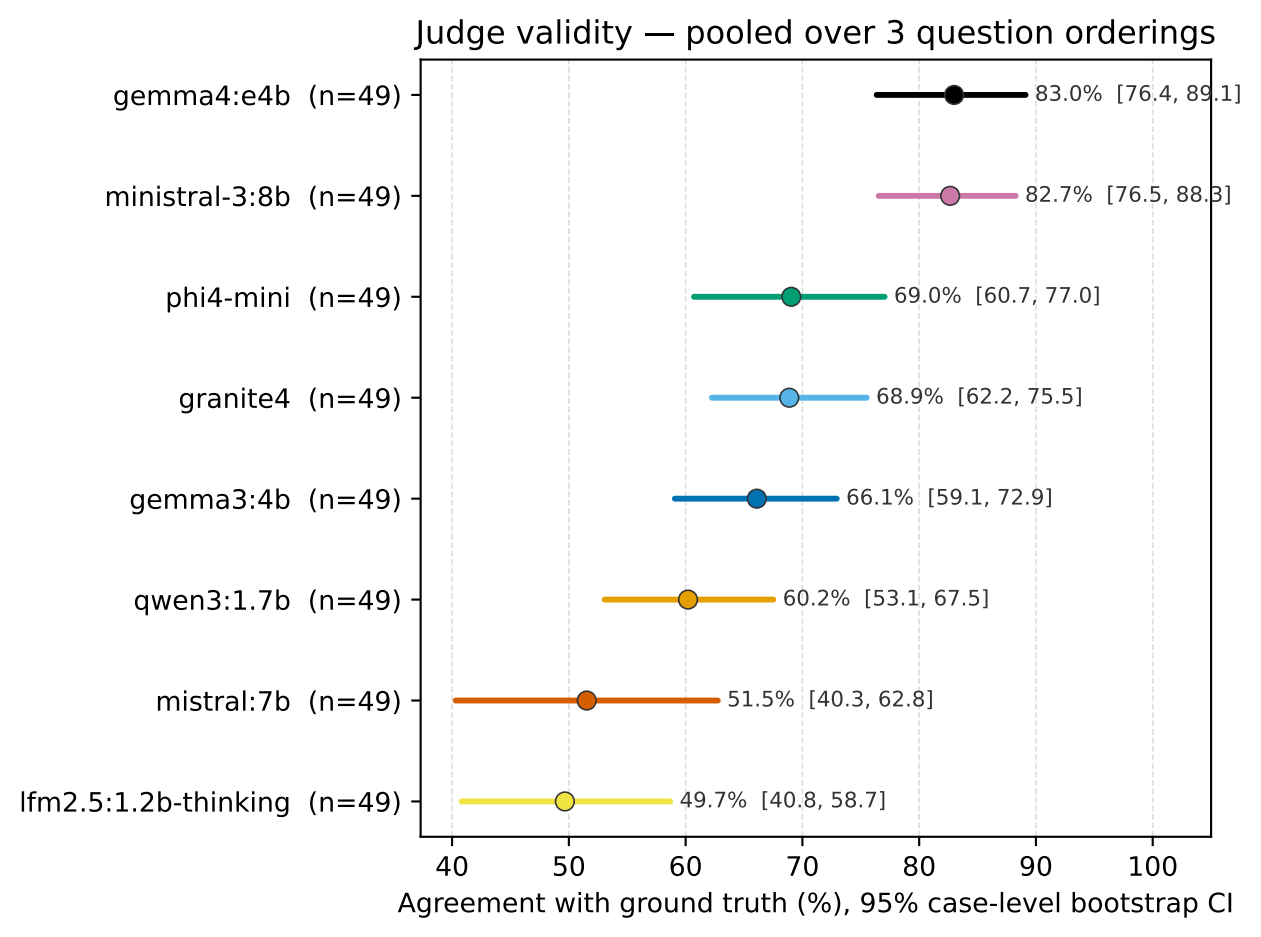
Judge validity — pooled over 3 question orderings
gemma4:e4b (n=49)
83.0% [76.4, 89.1]
ministral-3:8b (n=49)
82.7% [76.5, 88.3]
phi4-mini (n=49)
69.0% [60.7, 77.0]
granite4 (n=49)
68.9% [62.2, 75.5]
gemma3:4b (n=49)
66.1% [59.1, 72.9]
qwen3:1.7b (n=49)
60.2% [53.1, 67.5]
mistral:7b (n=49)
51.5% [40.3, 62.8]
lfm2.5:1.2b-thinking (n=49)
49.7% [40.8, 58.7]
40
50
60
70
80
90
100
Agreement with ground truth (%), 95% case-level bootstrap CI